

LECTURE 4

LANGUAGE MODELING AND RNNs PART 2

석사과정 장명준



KOREA
UNIVERSITY

DS·BA
Data Science &
Business Analytics

CONTENTS



1. VANISHING/EXPLODING GRADIENTS

2. LSTM & DEEP RNN

3. SCALING LARGE VOCABULARIES

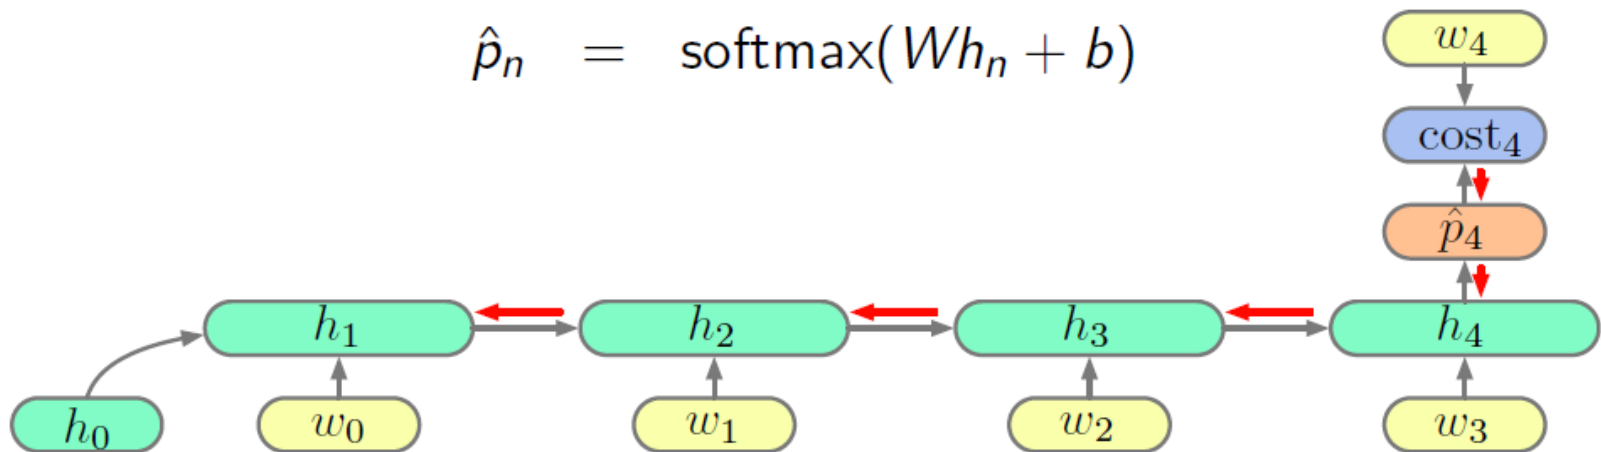
4. REGULARIZATION : DROPT OUT

VANISHING/EXPLODING GRADIENTS

Vanishing/Exploding gradients

$$h_n = g(V[x_n; h_{n-1}] + c)$$

$$\hat{p}_n = \text{softmax}(Wh_n + b)$$

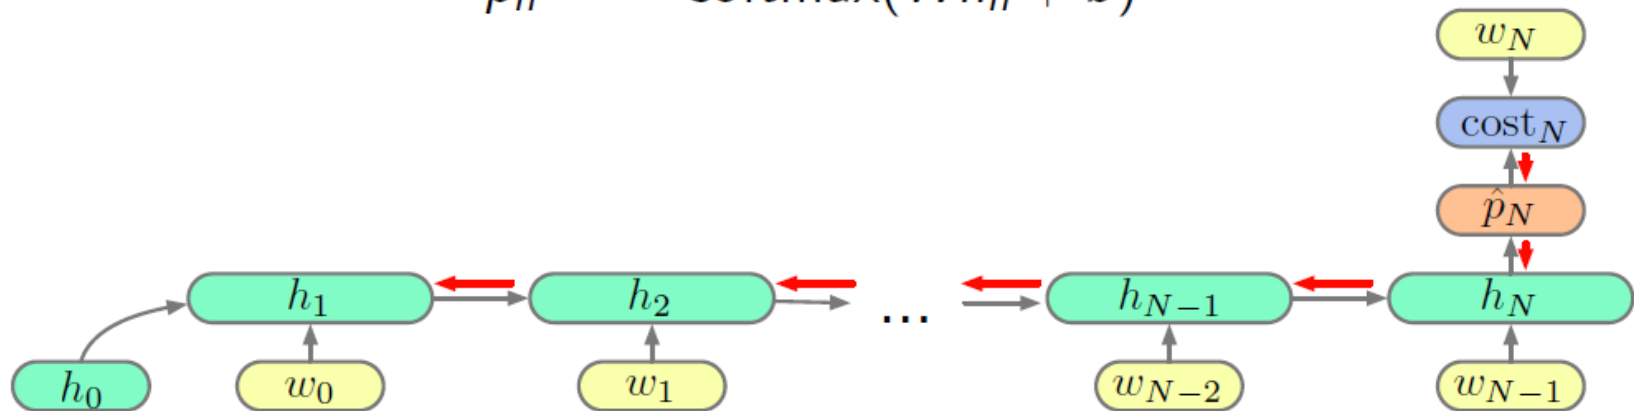


$$\frac{\partial \text{cost}_4}{\partial h_1} = \frac{\partial \text{cost}_4}{\partial \hat{p}_4} \frac{\partial \hat{p}_4}{\partial h_4} \frac{\partial h_4}{\partial h_3} \frac{\partial h_3}{\partial h_2} \frac{\partial h_2}{\partial h_1}$$

Vanishing/Exploding gradients

$$h_n = g(V[x_n; h_{n-1}] + c)$$

$$\hat{p}_n = \text{softmax}(Wh_n + b)$$



$$\frac{\partial \text{cost}_N}{\partial h_1} = \frac{\partial \text{cost}_N}{\partial \hat{p}_N} \frac{\partial \hat{p}_N}{\partial h_N} \left(\prod_{n \in \{N, \dots, 2\}} \frac{\partial h_n}{\partial h_{n-1}} \right)$$

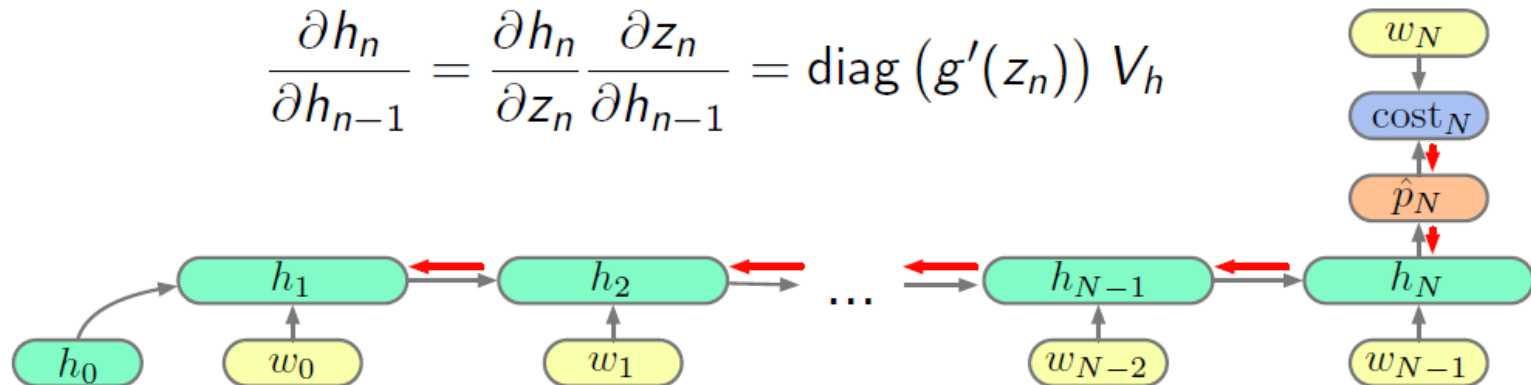
Vanishing/Exploding gradients

$$h_n = g(\underbrace{V_x x_n + V_h h_{n-1} + c}_{z_n}), \quad \frac{\partial \text{cost}_N}{\partial h_1} = \frac{\partial \text{cost}_N}{\partial \hat{p}_N} \frac{\partial \hat{p}_N}{\partial h_N} \left(\prod_{n \in \{N, \dots, 2\}} \frac{\partial h_n}{\partial z_n} \frac{\partial z_n}{\partial h_{n-1}} \right)$$

$$\frac{\partial h_n}{\partial z_n} = \text{diag}(g'(z_n))$$

$$\frac{\partial z_n}{\partial h_{n-1}} = V_h$$

$$\frac{\partial h_n}{\partial h_{n-1}} = \frac{\partial h_n}{\partial z_n} \frac{\partial z_n}{\partial h_{n-1}} = \text{diag}(g'(z_n)) V_h$$



Vanishing/Exploding gradients

$$\frac{\partial \text{cost}_N}{\partial h_1} = \frac{\partial \text{cost}_N}{\partial \hat{p}_N} \frac{\partial \hat{p}_N}{\partial h_N} \left(\prod_{n \in \{N, \dots, 2\}} \text{diag}(g'(z_n)) V_h \right)$$

$$\left\| \frac{\partial h_n}{\partial h_{n-1}} \right\| \leq \| \text{diag}(g'(z_n)) \| \| V_h \| \leq \beta_d \beta_h \quad \triangleright \quad \left\| \frac{\partial h_n}{\partial h_1} \right\| = \left\| \prod_{j=2}^n \frac{\partial h_j}{\partial h_{j-1}} \right\| \leq (\beta_d \beta_h)^{n-1}$$

If we use L2 norm

$$\left\| \frac{\partial h_n}{\partial h_1} \right\| = \left\| \prod_{j=2}^n \frac{\partial h_j}{\partial h_{j-1}} \right\| \leq \beta_d^{n-1} (\sqrt{\lambda_{\max}})^{n-1}$$

Since $\|A\|_2 = \max_{\|x\|_2=1} \|Ax\|_2 = \sqrt{\lambda_{\max}}$

If the largest eigenvalue of V_h is :

- 1, then gradient will propagate
- >1, exploding gradients
- <1, vanishing gradients

Vanishing/Exploding gradients

* Proof) $\|A\|_2 = \max_{\|x\|_2=1} \|Ax\|_2 = \sqrt{\lambda_{max}}$

Let's solve the problem below

$$\begin{aligned} & \text{Maximize}(\|Ax\|_2^2 = x^T A^T A x) \\ & \text{subject to } x^T x = 1 \end{aligned}$$

By using the Lagrange multiplier method

$$\frac{\partial x^T A^T A x}{\partial x} = \lambda \frac{\partial x^T x}{\partial x}$$

$$2A^T A x = 2\lambda x$$

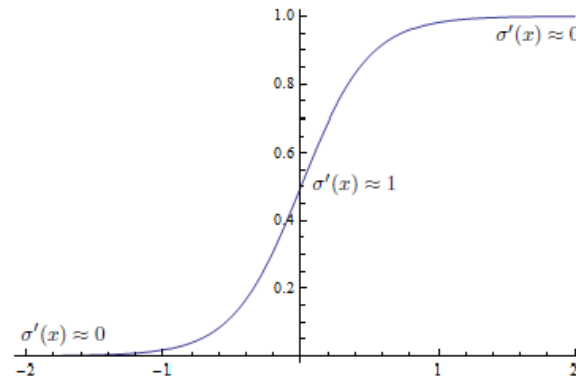
$$\therefore x^T A^T A x = \lambda x^T x = \lambda$$

It follows that

$$\|A\|_2 = \max_{\|x\|_2=1} \|Ax\|_2 = \left(\max_{x^T x=1} x^T A^T A x \right)^{1/2} = \sqrt{\lambda_{max}}$$

Vanishing/Exploding gradients

- Most of the times, the spectral radius (max eigenvalue) of V_h is small
- Also, many non-linearities($g(\cdot)$) can also shrink the gradient



Solutions:

1. Second order optimizers
2. Careful initialization
3. Changing the network architecture

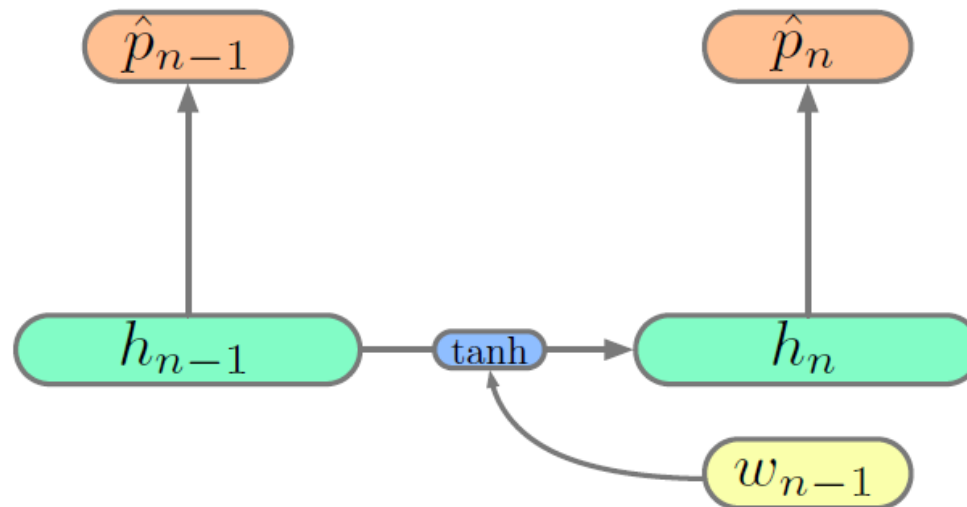
LSTM & DEEP RNN

Long Short Term Memory

Simple Recurrent Neural Network

$$h_n = \tanh(V[w_{n-1}; h_{n-1}] + c)$$

Update by multiplying



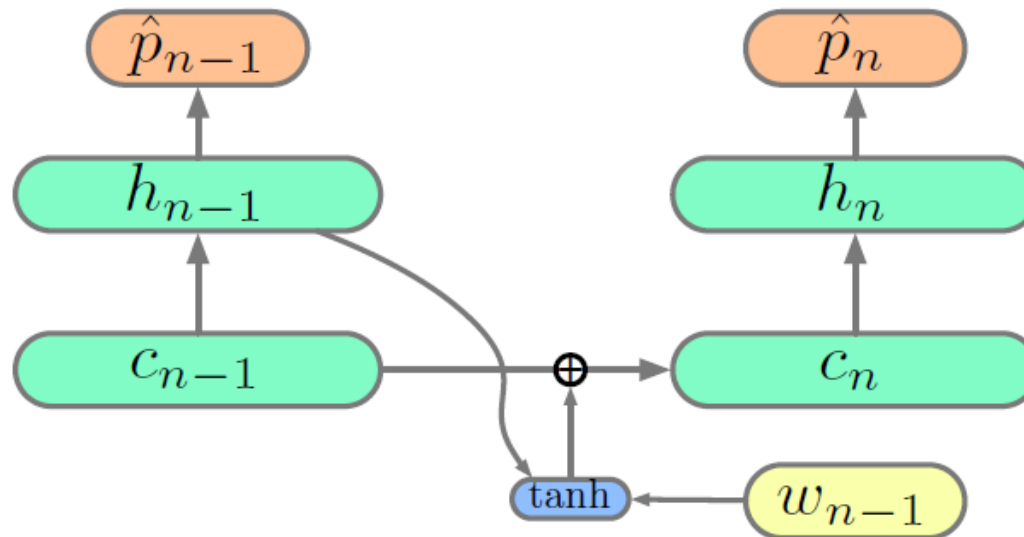
Long Short Term Memory

LSTM

- Introduce cell state : we can think this as memory

$$c_n = c_{n-1} + \tanh(V[w_{n-1}; h_{n-1}] + b_c)$$

Update by additive : gradient flow nicely



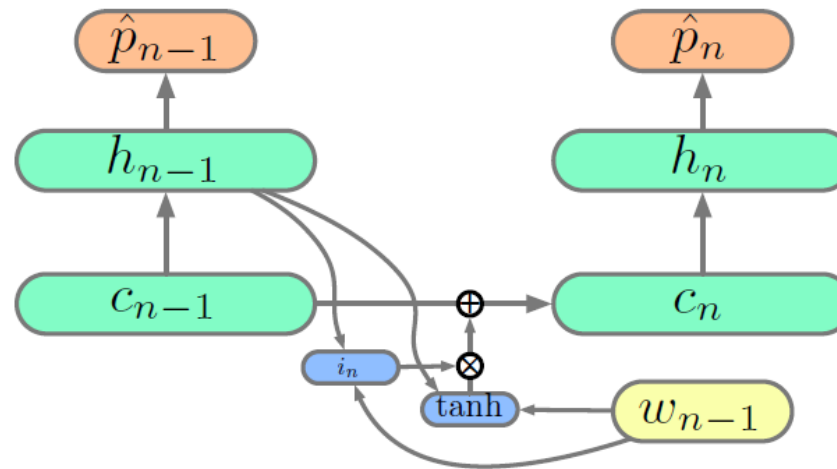
Need to balance : explode because of summation

Long Short Term Memory

LSTM

- Introduce gate for balancing : input gate

$$c_n = c_{n-1} + i_n \circ \tanh(V[w_{n-1}; h_{n-1}] + b_c)$$



$$i_n = \sigma(W_i[w_{n-1}; h_{t-1}] + b_i).$$

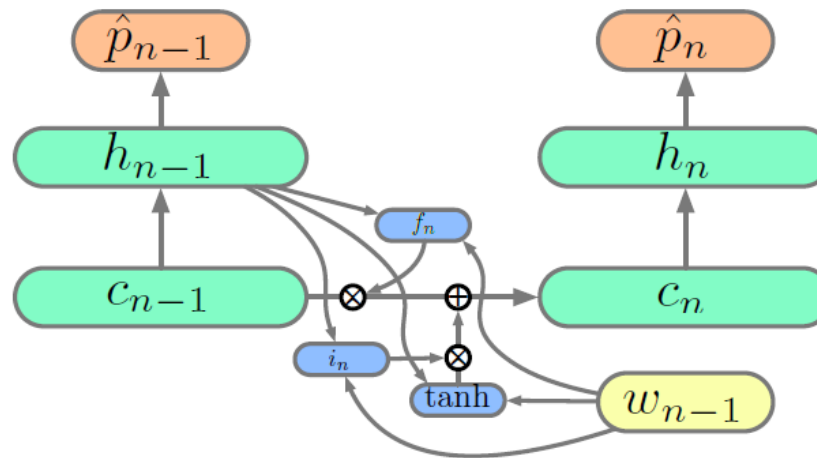
→ continuous vector : different gate for every hidden unit

Long Short Term Memory

LSTM

- gate for c_n : forget gate

$$c_n = f_n \circ c_{n-1} + i_n \circ \tanh(V[w_{n-1}; h_{n-1}] + b_c)$$



$$i_n = \sigma(W_i[w_{n-1}; h_{t-1}] + b_i),$$

$$f_n = \sigma(W_f[w_{n-1}; h_{t-1}] + b_f).$$

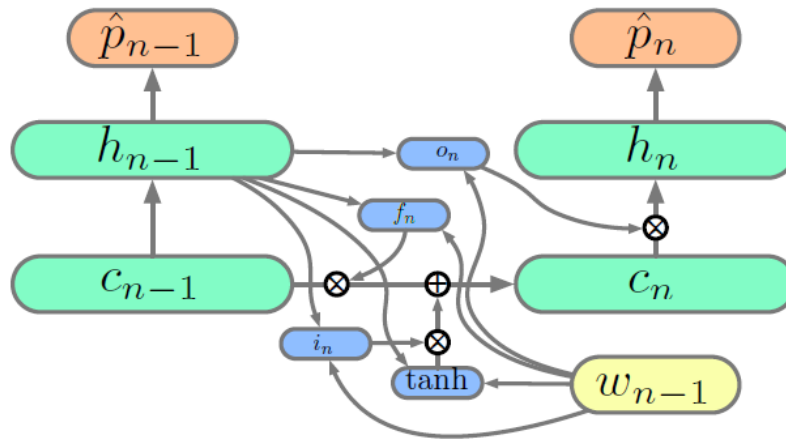
Long Short Term Memory

LSTM

- Why not put a gate on the output? : output gate

$$c_n = f_n \circ c_{n-1} + i_n \circ \tanh(V[w_{n-1}; h_{n-1}] + b_c)$$

$$h_n = o_n \circ \tanh(W_h c_n + b_h).$$



- Sigmoid : good to use as gate (0 ~ 1)
- tanh : can be replaced to other non-linearity

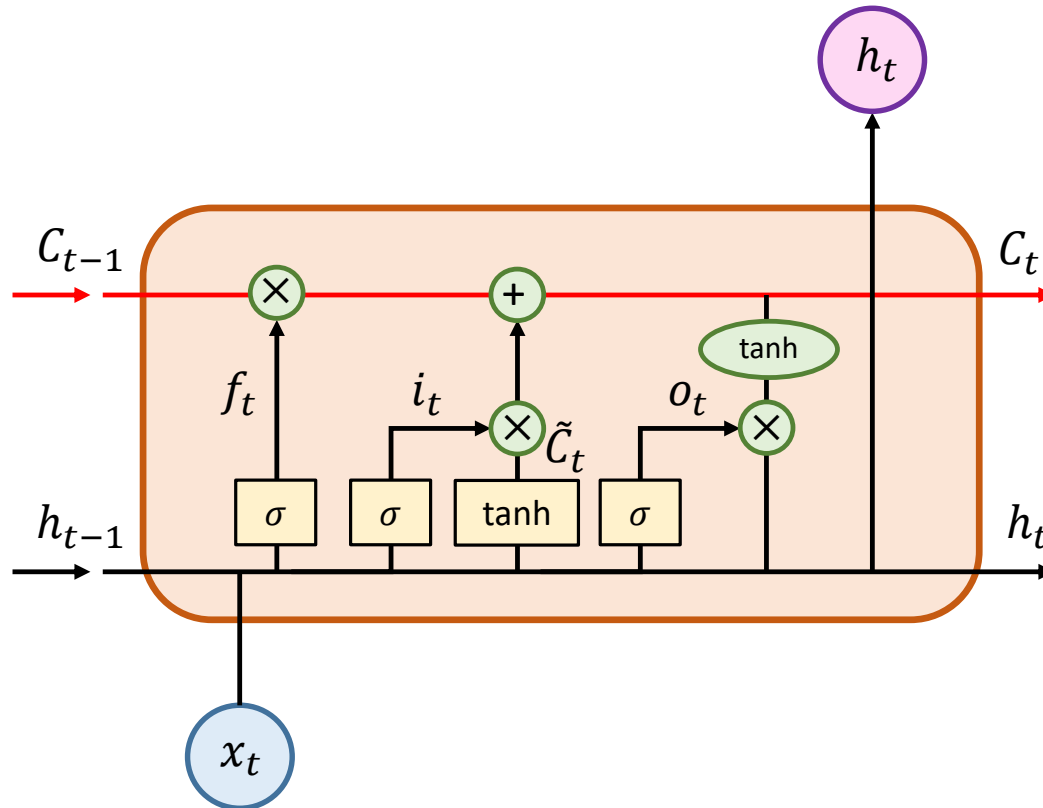
$$i_n = \sigma(W_i[w_{n-1}; h_{t-1}] + b_i),$$

$$f_n = \sigma(W_f[w_{n-1}; h_{t-1}] + b_f),$$

$$o_n = \sigma(W_o[w_{n-1}; h_{t-1}] + b_o).$$

Long Short Term Memory

LSTM



Long Short Term Memory

LSTM

$$c_n = f_n \circ c_{n-1} + (1 - f_n) \circ \hat{c}_n$$

$$\hat{c}_n = \tanh(W_c[w_{n-1}; h_{t-1} + b_c])$$

$$h_n = o_n \circ \tanh(W_h c_n + b_h)$$

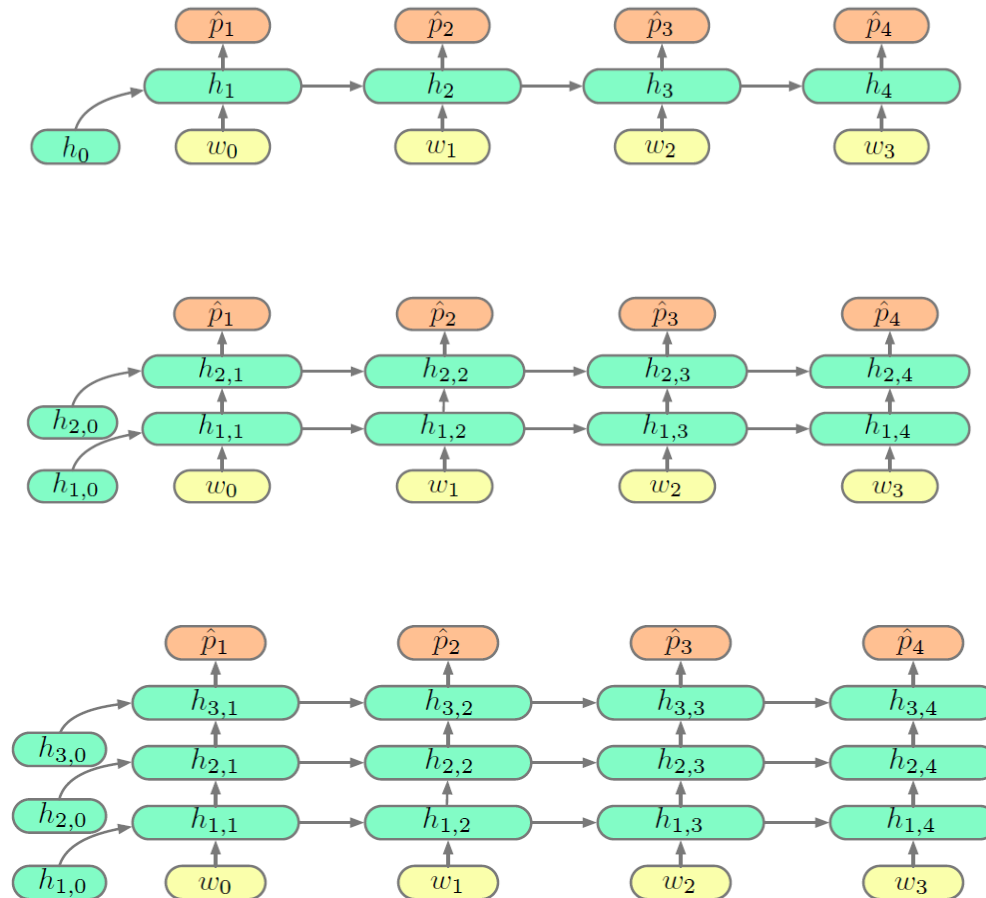
$$i_n = \sigma(W_i[w_{n-1}; h_{t-1} + b_i])$$

$$f_n = \sigma(W_f[w_{n-1}; h_{t-1} + b_f])$$

$$o_n = \sigma(W_o[w_{n-1}; h_{t-1} + b_o])$$

Long Short Term Memory

Deep RNN

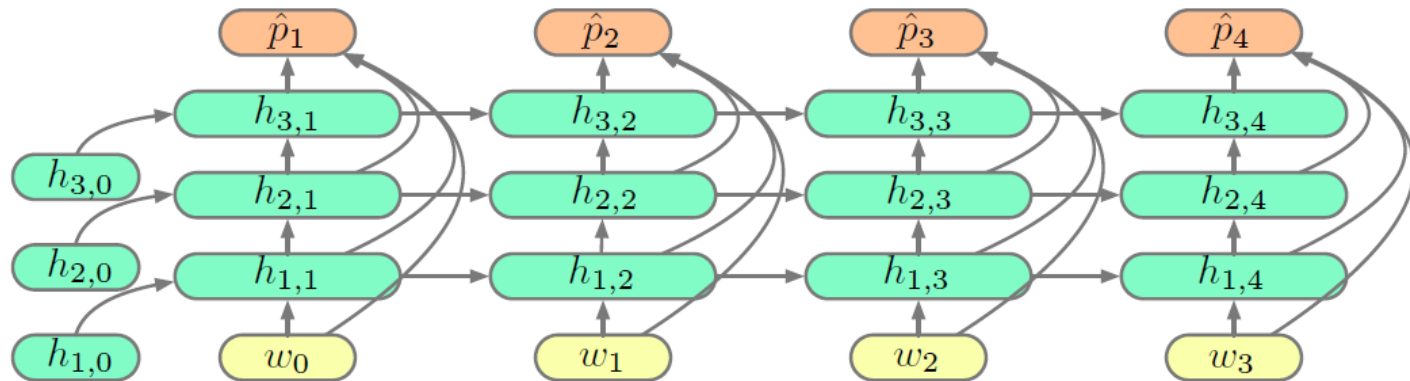


Deep RNN increases representational ability but also the memory

Long Short Term Memory

Deep RNN

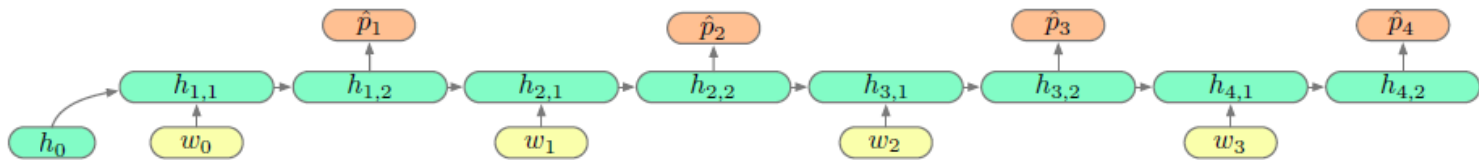
- Skip connection : Deep network is hard to back propagate



Long Short Term Memory

Deep RNN

- Increase in the time dimension
- This improves the representational ability
- But not the memory capacity : we still have 1 recurrent sequence



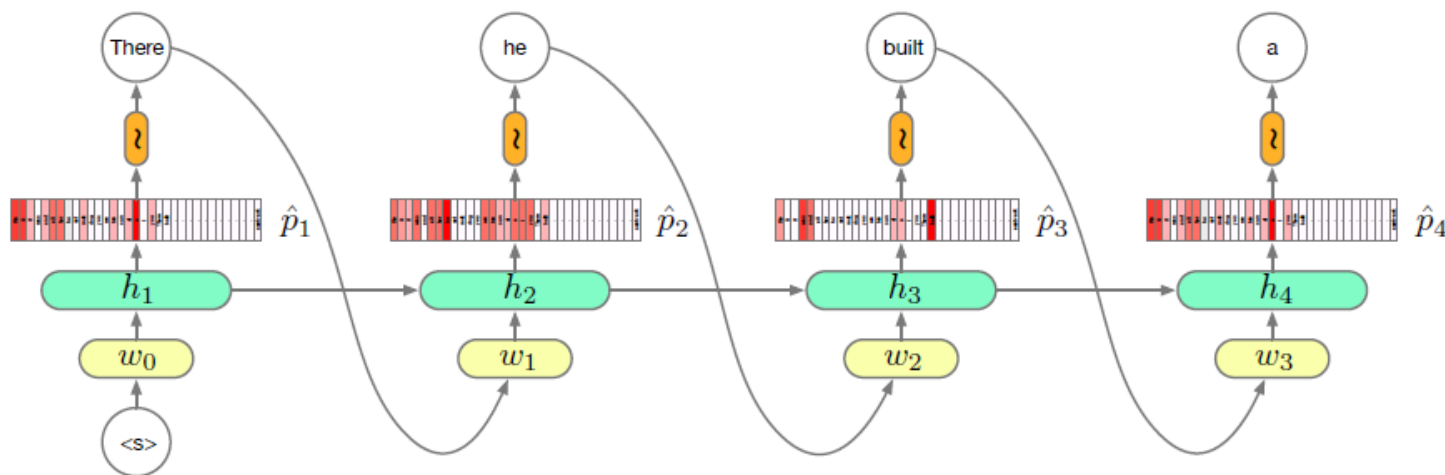
SCALING: LARGE VOCABULARIES

Scaling: Large Vocabularies

- Much of the computational cost is dominated by calculating:

$$\hat{p}_n = \text{softmax}(W_o h_n + b)$$

($|V| \times h$) \times ($h \times 1$)



Solutions:

1. **Short-list** : use the RNN model for the most frequent words, and traditional N-gram model for the rest, easy to use but lose RNN's main advantage (generalization to rare events)
2. **Batch local short-lists**

Scaling: Large Vocabularies

Batch local short-lists

- Approximate full partition function from a subset of the vocabulary

Probability

$$p(y_t | y_{<t}, x) = \frac{1}{Z} \exp(w_t^T \phi(y_{t-1}, z_t, c_t) + b_t) \quad \begin{aligned} z_t &= g(y_{t-1}, z_{t-1}, c_t) \\ c_t &= r(z_{t-1}, h_1, \dots, h_T) \end{aligned}$$

$$Z = \sum_{k: y_k \in V} \exp(w_k^T \phi(y_{t-1}, z_t, c_t) + b_k) \quad , \quad w_t^T: \text{target word vector}, b_t: \text{target word bias}$$

Calculate gradient

$$\nabla \log p(y_t | y_{<t}, x) = \nabla \varepsilon(y_t) - \sum_{k: y_k \in V} p(y_k | y_{<t}, x) \nabla \varepsilon(y_k) \quad \Rightarrow \quad E_P[\nabla \varepsilon(y)]$$

$\text{where } \varepsilon(y_j) = w_j^T \phi(y_{t-1}, z_t, c_t) + b_j$
 $\text{where } P \text{ denotes } p(y | y_{<t}, x)$

Main idea

- Approximate expectation by sampling with a small number of samples

Given a predefined proposal distribution Q and a set V' of samples from Q

$$E_P[\nabla \varepsilon(y)] \approx \sum_{k: y_k \in V'} \frac{w_k}{\sum_{k': y_{k'} \in V'} w_{k'}} \nabla \varepsilon(y_k)$$

$\text{where } w_k = \exp(\varepsilon(y_k) - \log Q(y_k))$

Scaling: Large Vocabularies

Batch local short-lists

In Practice

- Partition the training corpus and define a subset V' of the target vocabulary for each partition prior to training

Let V'_i refers to the subset of target words used for i^{th} partition and $Q_i(y_k) = \begin{cases} \frac{1}{|V'_i|} & \text{if } y_k \in V'_i \\ 0 & \text{otherwise} \end{cases}$

From equation

$$E_P[\nabla \varepsilon(y)] \approx \sum_{k: y_k \in V'} \frac{w_k}{\sum_{k': y_{k'} \in V'} w_{k'}} \nabla \varepsilon(y_k) \text{ where } w_k = \exp(\varepsilon(y_k) - \log Q(y_k))$$

$$\begin{aligned} p(y_t | y_{<t}, x) &\approx \frac{w_k}{\sum_{k': y_{k'} \in V'} w_{k'}} = \frac{\exp(\varepsilon(y_t) - \log Q(y_t))}{\sum_{k': y_{k'} \in V'} \exp(\varepsilon(y_{k'}) - \log Q(y_{k'}))} \\ &= \frac{\exp(\varepsilon(y_t) + |V'|)}{\sum_{k': y_{k'} \in V'} \exp(\varepsilon(y_{k'}) + |V'|)} \\ &= \frac{\exp(\varepsilon(y_t)) \exp(|V'|)}{\sum_{k': y_{k'} \in V'} \exp(\varepsilon(y_{k'})) \exp(|V'|)} \\ &= \frac{\exp(\mathbf{w}_t^T \phi(\mathbf{y}_{t-1}, \mathbf{z}_t, \mathbf{c}_t) + \mathbf{b}_t)}{\sum_{k': y_{k'} \in V'} \exp(\mathbf{w}_{k'}^T \phi(\mathbf{y}_{t-1}, \mathbf{z}_t, \mathbf{c}_t) + \mathbf{b}_{k'})} \end{aligned}$$

Scaling: Large Vocabularies

- Much of the computational cost is dominated by calculating:

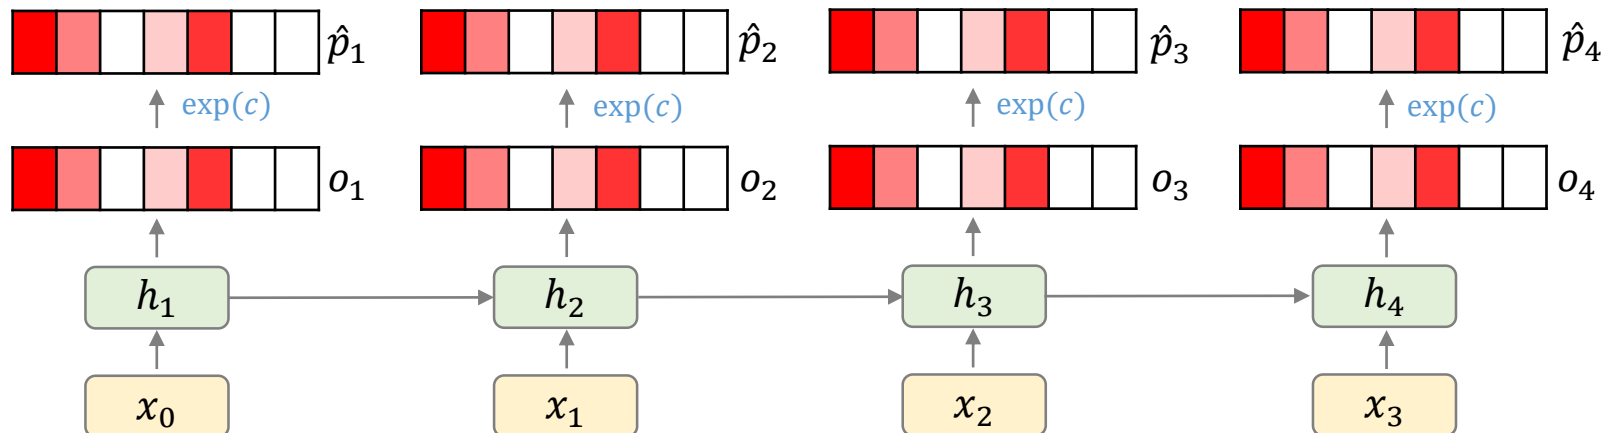
$$\hat{p}_n = \text{softmax}(W_o h_n + b)$$

Solutions:

3. Approximate the gradient/ change the objective :

consider maximizing likelihood by making the log partition function an independent parameter c

$$\hat{p}_n = \exp(W_o h_n + b) \times \exp(c)$$

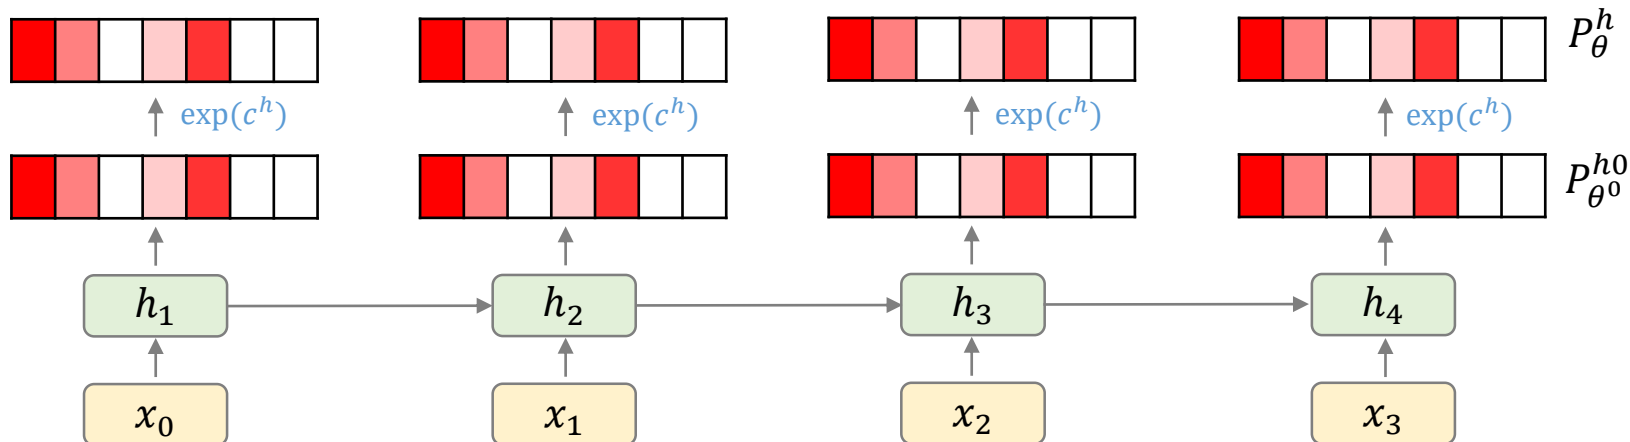


Leads to an ill defined objective

Scaling: Large Vocabularies

Noise-contrastive estimation

Reduce problem of softmax to that of **binary classification** discriminating of words that occur after a particular context h



$$P_{\theta}^h(w) = P_{\theta^0}^{h0}(w) \exp(c^h), \theta = \{\theta^0, c^h\}$$

$P_{\theta^0}^{h0}(w)$: unnormalized distribution, c^h : learned parameter

$P_d^h(w)$: empirical distribution of words that occur after a particular context h

$P_n(w)$: noise distribution (unigram)

Object

$$P_{\theta}^h(w) \approx P_d^h(w)$$

Scaling: Large Vocabularies

Noise-contrastive estimation

Assume that noise samples are k times more frequent than data samples

$$p^h(d, w) = \begin{cases} \frac{k}{1+k} \times P_n(w) & \text{if } d = 0 \\ \frac{1}{1+k} \times P_d^h(w) & \text{if } d = 1 \end{cases} \quad \rightarrow \quad \begin{aligned} p^h(D = 0|w) &= \frac{kP_n(w)}{P_d^h(w) + kP_n(w)} \\ p^h(D = 1|w) &= \frac{P_d^h(w)}{P_d^h(w) + kP_n(w)} \end{aligned}$$

Object
 $P_\theta^h(w) \approx P_d^h(w)$

$$\begin{aligned} p^h(D = 0|w) &= \frac{kP_n(w)}{P_d^h(w) + kP_n(w)} \\ p^h(D = 1|w) &= \frac{P_d^h(w)}{P_d^h(w) + kP_n(w)} \end{aligned} \quad \rightarrow \quad \begin{aligned} p^h(D = 0|w, \theta) &= \frac{kP_n(w)}{P_\theta^h(w) + kP_n(w)} \\ p^h(D = 1|w, \theta) &= \frac{P_\theta^h(w)}{P_\theta^h(w) + kP_n(w)} \end{aligned}$$

Scaling: Large Vocabularies

Noise-contrastive estimation

$$P^h(D = 0|w, \theta) = \frac{kP_n(w)}{P_\theta^h(w) + kP_n(w)}$$

$$P^h(D = 1|w, \theta) = \frac{P_\theta^h(w)}{P_\theta^h(w) + kP_n(w)}$$

Maximize the expectation of $\log P^h(D|w, \theta)$ under the mixture of the data and noise samples

$$J^h(\theta) = E_{P_d^h} \left[\log \frac{P_\theta^h(w)}{P_\theta^h(w) + kP_n(w)} \right] + kE_{P_n} \left[\frac{kP_n(w)}{P_\theta^h(w) + kP_n(w)} \right]$$

With the gradient

$$\begin{aligned} \frac{\partial}{\partial \theta} J^h(\theta) &= E_{P_d^h} \left[\frac{kP_n(w)}{P_\theta^h(w) + kP_n(w)} \frac{\partial}{\partial \theta} \log P_\theta^h(w) \right] + kE_{P_n} \left[\frac{P_\theta^h(w)}{P_\theta^h(w) + kP_n(w)} \frac{\partial}{\partial \theta} \log P_\theta^h(w) \right] \\ &= \sum_w \frac{kP_n(w)}{P_\theta^h(w) + kP_n(w)} \times (P_d^h(w) - P_\theta^h(w)) \frac{\partial}{\partial \theta} \log P_\theta^h(w) \end{aligned}$$

And that as $k \rightarrow \infty$

$$\frac{\partial}{\partial \theta} J^h(\theta) = \sum_w (P_d^h(w) - P_\theta^h(w)) \frac{\partial}{\partial \theta} \log P_\theta^h(w)$$

Scaling: Large Vocabularies

Noise-contrastive estimation

In practice, given a word w observed in context h , we compute its contribution to the gradient by generating k noise samples x_1, \dots, x_k and using the formula

$$\begin{aligned} \frac{\partial}{\partial \theta} J^h(\theta) &= E_{P_d^h} \left[\frac{k P_n(w)}{P_\theta^h(w) + k P_n(w)} \frac{\partial}{\partial \theta} \log P_\theta^h(w) \right] + k E_{P_n} \left[\frac{P_\theta^h(w)}{P_\theta^h(w) + k P_n(w)} \frac{\partial}{\partial \theta} \log P_\theta^h(w) \right] \\ &= \frac{k P_n(w)}{P_\theta^h(w) + k P_n(w)} \frac{\partial}{\partial \theta} \log P_\theta^h(w) + k \sum_{i=1}^k \frac{1}{k} \frac{P_\theta^h(x_i)}{P_\theta^h(x_i) + k P_n(x_i)} \frac{\partial}{\partial \theta} \log P_\theta^h(x_i) \end{aligned}$$

Since distributions for different contexts share parameters, we cannot learn these distributions independently

$$J(\theta) = \sum_h P(h) J^h(\theta)$$

$P(h)$: empirical context probabilities

Scaling: Large Vocabularies

- Much of the computational cost is dominated by calculating:

$$\hat{p}_n = \text{softmax}(W_o h_n + b)$$

Solutions:

4. Factorize the output vocabulary :

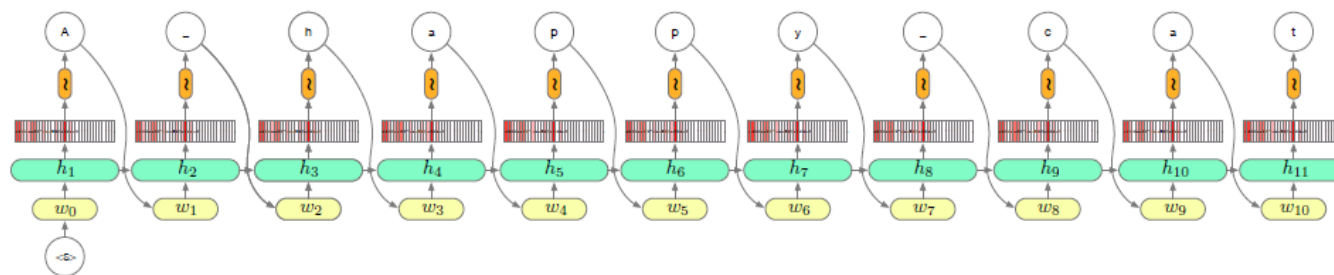
- Make assumption about distribution
- There are lots of algorithms for grouping vocabularies into classes

$$p(w_n | \hat{p}_n^{class}, \hat{p}_n^{word}) = p(class(w_n) | \hat{p}_n^{class}) \times p(w_n | class(w_n), \hat{p}_n^{word})$$

- Assuming balanced classes, this gives a \sqrt{V} speed up
- Ex) $V = 1,000,000$, classes = 1,000

Sub-Word Level Language Modeling

- Result in much smaller softmax and no unknown words
- Model can capture subword structure and morphology : disunited, disinherited, disinterested
- Longer sequences, hard to back propagate
- A lot of structure in languages, bottom low is word and we want to learn correlation between words
- Characters cannot learn correlation between the words



REGULARIZATION

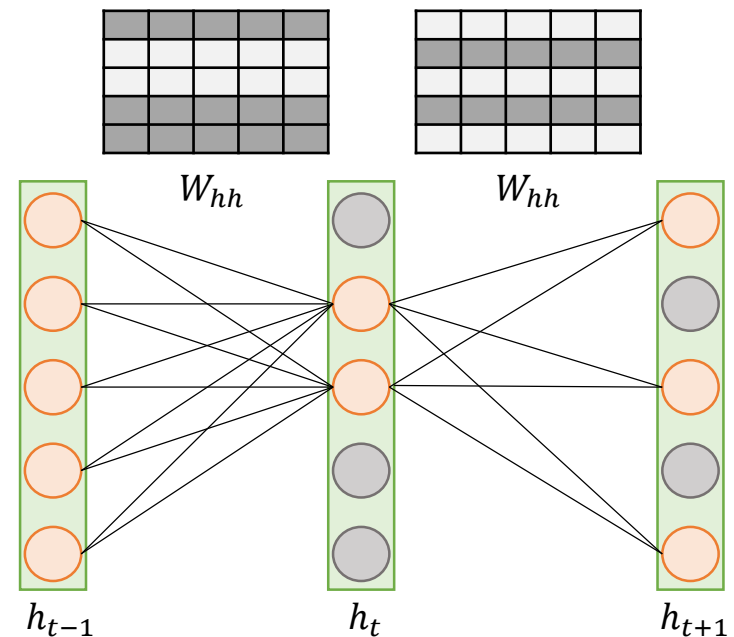
Regularization: Drop out

- Dropout is ineffective when applied to recurrent connections
- If we repeat couple of times, in expectation, every hidden unit will be dropped out

$$h_t = W_{hh} h_{t-1}$$

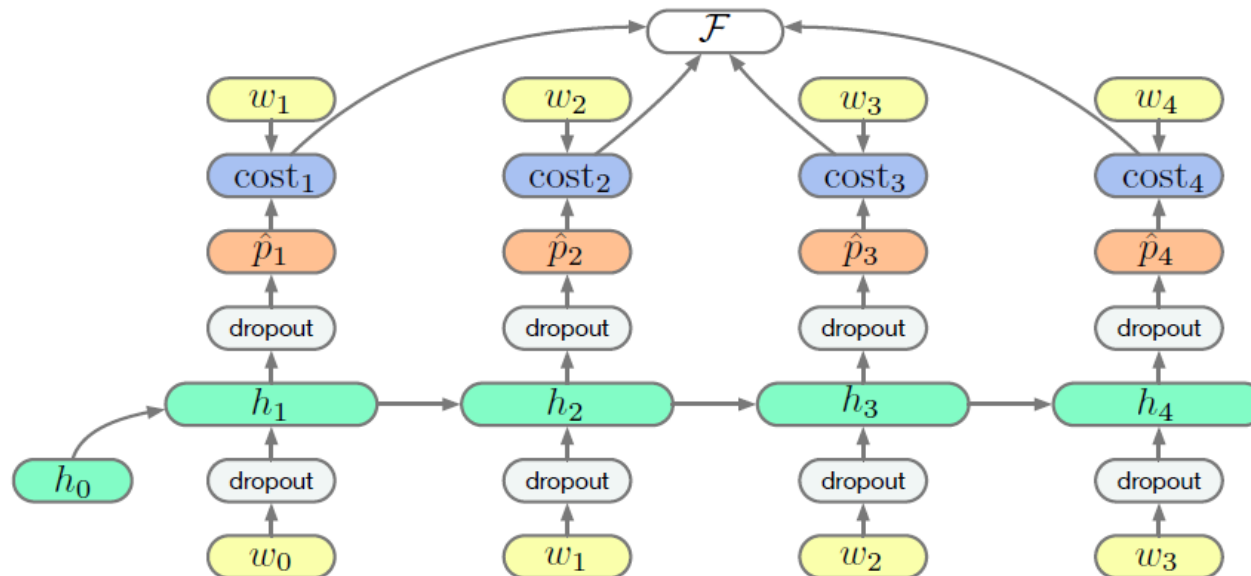
$h \times 1$ $h \times h$ $h \times 1$

$$\begin{bmatrix} h_{t,1} \\ h_{t,2} \\ h_{t,3} \\ \vdots \\ h_{t,h} \end{bmatrix} = \begin{bmatrix} w_{1,1} & w_{1,2} & \dots & w_{1,h} \\ w_{2,1} & w_{2,2} & \dots & w_{2,h} \\ \vdots & \dots & \ddots & \vdots \\ w_{h,1} & w_{h,2} & \dots & w_{h,h} \end{bmatrix} \begin{bmatrix} h_{t-1,1} \\ h_{t-1,2} \\ h_{t-1,3} \\ \vdots \\ h_{t-1,h} \end{bmatrix}$$



Regularization: Drop out

- Dropout is ineffective when applied to recurrent connections
- If we repeat couple of times, in expectation, every hidden unit will be dropped out
- Instead, put drop out everywhere else in the network (Input, Output layer)



- But recurrent connections are not regularized

